

Prompts Used During the AI Review

FNDR Project Design — Submission Document 2

Prompts only; design-document text and AI responses omitted

Initial Prompt

Review our complete FNDR project design document as an additional senior software/AI systems design reviewer. Critically evaluate the overall architecture, completeness, component interfaces, data models and data flow, privacy and security, reliability, performance, scalability, implementation complexity, inconsistencies, missing requirements, and risks before Alpha development.

Pay particular attention to whether the formal requirements and architecture remain consistent with the later project timeline, including semantic memory, computer-use agents, VLM grounding, multi-step reasoning, and task execution. Do not assume every design decision is correct merely because it appears in the document. Identify architectural contradictions and areas that need additional design work.

Main Subsequent Prompt 1

Reconsider the design assuming FNDR is evolving from primarily semantic search into a memory-to-action system. The semantic memory and retrieval system should provide relevant historical context to an agent that can select reusable skills and perform constrained browser/desktop workflows. The agent should use an external ChatGPT model primarily for high-level planning while retrieval and execution infrastructure remain part of FNDR.

Evaluate whether the current architecture actually supports this direction and identify what new modules, interfaces, data models, and boundaries would be required.

Main Subsequent Prompt 2

Examine FNDR's proposed use of reinforcement learning and evaluations more critically. The current intent is not to update the foundation model's weights. Instead, the team plans to run benchmark tasks and use failures to improve skill Markdown, prompts, orchestration, tool choice, and context retrieval.

Determine whether “reinforcement learning” is technically accurate terminology and propose an evaluation architecture that separately measures retrieval, planning, skill selection/orchestration, action selection, and end-to-end task completion.

Main Subsequent Prompt 3

Review the privacy and safety implications of combining FNDR's local-first memory architecture with an external ChatGPT planner. Assume raw screenshots remain transient and local, but selected retrieved context and metadata may be provided to the model. Sensitive categories should be filtered before transmission. Agent workflows should be allowlisted, and early versions should require user confirmation for actions.

Identify contradictions with the existing privacy claims and recommend a clear local/cloud trust boundary and action-permission model.

Main Subsequent Prompt 4

Re-evaluate the project's Alpha scope. The team wants Alpha to emphasize high-quality memory retrieval and correct skill selection while still demonstrating at least one compelling multi-step end-to-end workflow. Evaluation results should drive improvements to versioned skills, prompts, context selection, and orchestration.

Identify functionality in the current roadmap that is insufficiently specified, unnecessarily broad, or lower priority compared with this memory-to-action steel thread. Produce the final review using the five most significant and actionable findings.