Team Evaluation of Five AI Review Findings

FNDR Project Design — Submission Document 3

Team FNDR

AI Tool, Model, and Review Process

The team provided the complete FNDR Final Design Document to ChatGPT using GPT-5.6 Sol. Rather than relying on a single response, the team iteratively clarified the intended product direction, agent architecture, privacy model, skill system, evaluation strategy, and Alpha goals. The AI was then asked to reconsider its findings using those constraints and produce a final critical review. The team selected the five findings below as the most significant and actionable.

Finding 1 — Agent architecture is missing from the formal system design

Team decision	**ACCEPT**

Summary: The existing seven-module architecture primarily describes semantic capture and retrieval, while agent planning, skills, orchestration, execution, and evaluation appear mainly in the later timeline.

Why: The team agrees that the design now needs to represent the memory-to-action path explicitly rather than treating the agent as an implementation detail.

Resulting change: Add an explicit flow connecting retrieval to context assembly, a ChatGPT-based planner, skill/tool orchestration, constrained execution, verification, and evaluation. Skills will be reusable versioned definitions containing instructions, tools, constraints, failure handling, and examples.

Finding 2 — Cloud-assisted reasoning conflicts with the current privacy language

Team decision	**PARTIALLY ACCEPT**

Summary: The current document describes normal FNDR operation as fully local with no outbound model requests, while the newer agent direction may send selected retrieved context and metadata to a cloud planner.

Why: The team agrees that this changes the privacy boundary, but does not agree that FNDR should abandon its local-first architecture. The core memory system should remain local and cloud-assisted agent functionality should be treated as experimental and distinct.

Resulting change: Explicitly separate local memory operation from experimental cloud-assisted agent mode. Sensitive categories will be filtered before context is sent externally, and only selected retrieval results and necessary metadata will be provided to the planner.

Finding 3 — Agent evaluation is insufficiently defined

Team decision	**ACCEPT**

Summary: The current QA section focuses on the original semantic-memory pipeline and does not adequately measure the newer agent system.

Why: The team agrees that the current work is more accurately described as evaluation-driven improvement than foundation-model reinforcement learning because the team is not currently updating model weights.

Resulting change: Create separate evaluations for retrieval quality, planning, skill/orchestration selection, action selection, and end-to-end completion. Benchmark failures will drive improvements to prompts, skills, orchestration, and retrieval strategies, with results compared between versions.

Finding 4 — Agent safety and permission boundaries need to be explicit

Team decision	**ACCEPT**

Summary: The timeline mentions safety testing and human-in-the-loop gates, but the execution permissions and confirmation model are not formally defined.

Why: The team agrees that unrestricted GUI control would introduce unnecessary risk while the agent and evaluation harness are still being developed.

Resulting change: Alpha workflows will use an allowlisted set of actions and skills. User confirmation will initially be required before agent actions are executed. The design may later support more granular skill/workflow permissions as reliability is demonstrated through evaluations.

Finding 5 — The roadmap contains too many disconnected advanced systems

Team decision	**ACCEPT**

Summary: The roadmap includes meeting intelligence, GraphStore, knowledge-graph visualization, cloud/mobile infrastructure, and other advanced systems that could compete with the central memory-to-action objective.

Why: The team agrees that Alpha should validate the core product hypothesis before distributing effort across unrelated advanced features.

Resulting change: Prioritize the memory-to-action steel thread: retrieve useful historical context → select the correct skill → execute a constrained workflow. The team still intends to demonstrate at least one substantial multi-step end-to-end workflow, while unrelated advanced features remain lower priority until this path is reliable.